

XX

by

Rustam Issabekov



Thesis

Submitted in total fulfillment of the
requirements for the degree of
Doctor of Philosophy

Principal supervisor: XXXXX

Associate supervisor: XXXX

Associate supervisor: XXXX

**School of Science, Information Technology and
Engineering
University of Ballarat**

PO Box 663

University Drive, Mount Helen
Ballarat, Victoria 3353, Australia

February, 2014

Abstract

Statement of Authorship

This thesis contains no work extracted in whole or in part from a thesis, dissertation or research paper previously presented for another degree or diploma except where explicit reference is made. No other persons work has been relied upon or used without due acknowledgment in the main text and bibliography of the thesis.

Rustam Issabekov
February 18, 2014

Acknowledgments

Rustam Issabekov

University of Ballarat

February 2014

Contents

Abstract	ii
Statement of Authorship	iii
Acknowledgments	iv
List of Tables	vii
List of Figures	viii
List of Publications	x
1 Literature review.	1
1.1 Elements of reinforcement learning	1
1.1.1 The Problem.	1
1.1.2 Agent and environment	2
1.1.3 Agent’s goals.	4
1.1.4 Dynamic Programming influence.	5
1.1.5 Q-Learning.	9
1.1.6 Off-policy and On-policy algorithms.	10
1.1.7 S.A.R.S.A.	12
1.2 Multi-objective Optimization.	12
1.2.1 Multi-objective Problem.	13
1.2.2 Tradeoffs.	13

1.2.3	Dominance.	13
1.2.4	The Pareto Front.	14
1.2.5	Selecting a solution in the Pareto Front.	14
1.3	Multi-objective Reinforcement Learning Research.	17
1.3.1	Linear Temporal Difference Learning.	17
1.3.2	Non-linear Temporal Difference Learning.	20
1.3.3	Simultaneous Learning of More Than One Policy.	23
Introduction		1
2 Empirical Evaluation.		27
2.1	Introduction.	27
2.2	The Details of the Empirical Evaluation Study.	29
2.2.1	Performance metrics.	30
2.2.2	Experiment Structure.	32
2.2.3	Benchmarking Software.	37
2.3	Benchmarks.	40
2.3.1	Deep sea treasure.	41
2.3.2	MO-Puddleworld.	43
2.3.3	MO-Mountain Car.	44
2.3.4	Resource gathering.	44
Conclusions		46
References		46

List of Tables

2.1 **Left-hand side table** is the Q-function for the state s_{t+1} for the policy $\pi_{\vec{w}_1}$. **Right-hand side table** is the Q-function for the state s_{t+1} for the policy $\pi_{\vec{w}_2}$ 37

List of Figures

1.1	The agent-environment cycle (Sutton and Barto, 1998).	2
1.2	The Pareto front is represented with black points. Each black point is a non-dominated solution. Grey points represent dominated solutions. Each gray point is dominated by at least one black point.	15
2.1	The figure shows how points from a Pareto front approximation and a reference point together form the region of objective space. The black dots are Pareto optimal solutions and the red dot is the reference point.	31
2.2	Blue circles represent states; Black arrows represent the actions. In every state left and right actions are available. State transition is deterministic and an action always leads to the same next state. Next to every arrow there is a vector of immediate rewards for choosing this action. The overall return starting from starting state is shown at the bottom of each branch.	32
2.3	Each blue dot is a unique policy from the example MDP problem. NOTE: Because the example MDP does not have dominated policies all four available policies together constitute the Pareto front for this problem.	33
2.4	The blue dot which represents the policy with return [60,60] is marked with red circle around. The single policy algorithm mentioned above will learn the Q-function for the policy associated with the marked blue dot.	34

2.5	The blue dot which represents the policy with return $[60,60]$ is marked with red circle around. The single policy algorithm mentioned above will learn the Q-function for the policy associated with the marked blue dot.	36
2.6	The RL-Glue standard.	39
2.7	Grid world for Deep Sea Treasure problem.	42
2.8	Pareto front for Deep Sea Treasure problem.	42
2.9	Grid world for the Puddleworld problem.	43
2.10	Pareto front for the Puddleworld problem.	44
2.11	Grid world for the Resource Gathering problem.	45
2.12	Optimal Policies for the Resource Gathering problem.	45

List of Publications

Chapter 1

Literature review.

The purpose of this review is to summarise the literature on currently available multi-objective reinforcement learning algorithms. Section 1 of this chapter is an introduction to single objective reinforcement learning including common terms, definitions and notation. The second section is devoted to the multi-objective optimization and its influence on multi-objective reinforcement learning. Finally, Section 3 is a literature review of the important papers in multi-objective reinforcement learning.

1.1 Elements of reinforcement learning

Multi-objective reinforcement learning is a direct extension from single objective reinforcement learning. Both these disciplines share much of the same notation and terminology.

This section will introduce the core concepts of reinforcement learning in the context of a single objective function. The extension of these concepts to multiple objectives will be addressed in section 1.3.

1.1.1 The Problem.

Consider a multi-step decision making process with unknown reward and state transition functions. We can look at this process as an optimization problem in which the aim

is to make the correct decision at each step in order to maximize some measure of the rewards received. Any method that solves this optimization problem can be defined as a reinforcement learning method.

1.1.2 Agent and environment

A reinforcement learning problem can be further divided into two parts: an agent and an environment. The agent or so called decision maker or a learning algorithm is responsible for actions at each time step. The actions chosen by the agent affect the environment and cause it to transition from one state to another. This dynamic interaction constitutes the agent-environment cycle shown in Figure 1.1.




Figure 1.1: The agent-environment cycle (Sutton and Barto, 1998).

The agent interacts with the environment with the help of three signals: a state signal, which describes the current state of the environment; an action signal, by which the agent interacts with the environment and a scalar reward signal, a feedback from the environment on the desirability of the immediate consequences of the action. At each time step of the agent-environment cycle, the agent senses the current state of the environment and applies an action. This action causes the environment to transition into a new state. After each state transition a reward (possibly zero) is passed to the agent. This reward evaluates the quality of the taken action. Along with the reward the agent senses the new environmental state. This process happens at each time step. The agent does not control

state transitions and reward received. Instead the transition function and a reward function along with the set of possible actions and the set of all possible states belong to the environment which is often assumed to be a Markov decision process (MDP) (Bertsekas, 1995).

Mathematically speaking a MDP consists of 2 sets and 2 functions. The two sets are the set of all possible states S and a set of all possible actions A . The two functions are the transition function f and the reward function ρ . During interaction between the agent and the environment transition function f changes the state of the environment in response to the action signal. Right after state transition the reward function ρ produces a scalar reward signal, i.e. immediate reward for an agent. During a time step t the agent applies action a_t in the state s_t . The state changes to s_{t+1} according to the transition function $f : S \times A \rightarrow S$:

$$s_{t+1} = f(s_t, a_t)$$

And the agent receives the scalar reward r_{t+1} , according to the reward function $\rho : S \times A \rightarrow \mathbb{R}$:

$$r_{t+1} = \rho(s_t, a_t)$$

It is important to notice that both functions f and ρ can be either deterministic or stochastic.

The agent chooses actions according to its policy. The policy manifests which actions to pick at particular states $\pi : S \rightarrow A$:

$$a_t = \pi(s_t)$$

Because environment is represented as a Markov decision process, it has the Markov property. That is given only ρ, f, a_t, s_t it is possible to determine the probability distributions of s_{t+1} and r_{t+1} . This implies that all of the information required to make optimal decisions is contained within the current state vector s_t - the agent need not take into account any historical information about previous states and actions. A good example is the chess game. One doesn't need to know the previous moves to decide the current move. It is enough to know the current positions of all figures on a board to make a decision.

1.1.3 Agent's goals.

On every stage of the decision process an agent chooses an action. The quality of the chosen action is evaluated and a reward is returned to the agent. This setting allows us to formalize the goal of the agent. The goal of the agent is an optimization problem of maximizing the amount of rewards throughout a course of interaction between agent and environment. In other words the agent must decide which sequence of actions will result in maximum of the rewards.

This formalization of the goal naturally reflects the learning process found in people and animals and thus allows to model many of the problems found in real life.

A *return* is a cumulative sum of rewards obtained by the agent starting at some state s_0

$$\gamma^0 r_1 + \gamma^1 r_2 + \gamma^2 r_3 + \dots$$

Here, the discount factor $\gamma \in [0,1)$. First, this γ guarantees convergence over the infinite horizon and second it denotes that the agent gives more weight to immediate rewards rather than delayed rewards.

Alternatively the return could be written like this:

$$R^\pi(s_0) = \sum_{t=0}^{\infty} \gamma^t r_{t+1} = \sum_{t=0}^{\infty} \gamma^t \rho(s_t, \pi(s_t)).$$

And again γ is the discount factor and $s_{t+1} = f(s_t, \pi(s_t))$ for $t \geq 0$ and π is the policy being used.

Essentially, γ converts an infinite horizon MDP into a finite one. It is important to notice that in some cases it is possible to remove γ completely. For example, consider a case of a finite MDP with a -1 penalty for each move of an agent. In this setting this penalty will essentially work as γ .

1.1.4 Dynamic Programming influence.

Dynamic programming (Bellman, 1956) is the scientific discipline centered around calculating the optimal policies for a given MDP. It has a strong assumption that we have a perfect model of the MDP (full knowledge of the transition and the reward functions).

Reinforcement learning has a strong connection to the dynamic programming. As was said before the dynamic programming requires the knowledge of both transition and reward functions. But in some scenarios this information is not available to us. Cases like these are the natural application of the reinforcement learning. The following paragraphs will describe terms that are common to both the dynamic programming and the reinforcement learning.

Value Functions and Bellman Equation.

A large optimization problem of maximizing a reward over all stages can be broken into a sequence of smaller problems. Each of these small problems is centered around picking an action which will maximize the combination of the immediate reward and a cumulative sum of rewards that follows the next state of the environment. This requires us to have some means to store and access the cumulative reward associated with each state of an environment.

A *Value function* is a function that takes a state (or a combination of a state and an action) and returns the expected future reward that follows that state (or state-action).

In essence, a value function describes the "goodness" or value of visiting a state and allows the agent to compare states with each other.

The value of a state s_t , at time step t , depends on the rewards r_{t+1}, r_{t+2}, \dots . Which in turn depend on actions being taken at time steps $t, t+1, \dots$. Actions of course depend on a policy being used by a decision-maker. Thus a value function of a state s is denoted $V^\pi(s)$, with regard to some policy π , which underlies a decision making process.

Formally we can write

$$V^\pi(s) = E_\pi\{R_t \mid s_t = s\} = E_\pi\left\{\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid s_t = s\right\}.$$

Function V^π is called the *state-value function* for policy π . The state-value functions are good for conveying the theory of multi-step decision processes. In practice, however, they are replaced by so called *action-value functions*. The action-value function estimates the value of not only visiting a state s but also of taking some action a at that state, and subsequently following a policy π . Formally we can write

$$Q^\pi(s, a) = E_\pi\{R_t \mid s_t = s, a_t = a\} = E_\pi\left\{\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid s_t = s, a_t = a\right\}.$$

We can further notice that

$$\begin{aligned} V^\pi(s) &= E_\pi\{R_t \mid s_t = s\} \\ &= E_\pi\left\{\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid s_t = s\right\} \\ &= E_\pi\left\{r_{t+1} + \gamma \sum_{k=0}^{\infty} \gamma^k r_{t+k+2} \mid s_t = s\right\} \\ &= E_\pi\{r_{t+1} + \gamma V^\pi(s_{t+1}) \mid s_t = s\}. \end{aligned}$$

This equation is called the *Bellman equation* for V^π . The Bellman equation describes the connection between the value of a current state and the values of possible successor

states. This fundamental recursive relationship forms the foundation for a number of techniques to estimate correct values for V^π .

Optimal Value Functions.

Value functions allow us to compare different policies with each other, and naturally (because we are working with finite MDPs) there will be at least one best policy. To illustrate this we can visualize the given MDP as a tree. The tree's root is a starting state. Each branch of the tree represents a sequence of decisions (a policy) and the corresponding rewards. When the tree is finite there are a finite sequence of scalar rewards corresponding to each branch of the tree. One of those sequences will be the maximum. We can still apply the same logic even if the given MDP is infinite. In fact this is one of the reasons for using the discounting factor in the definition of return. The discounting factor allows us to represent the infinite MDP as the finite tree of decisions. We define a policy π to be better than a policy π' if it achieves higher or equal expected return for all states $s \in S$. Designation for optimal policy is π^* .

Any policy leads to a value function and naturally the optimal policy π^* leads to the optimal value function V^* . Formally we define V^* as

$$V^*(s) = \max_{\pi} V^\pi(s), \quad \text{for all } s \in S.$$

Now we can reformulate the reinforcement learning problem as finding the optimal policy π^* and associated optimal value function V^* .

Value Iteration.

We now turn to how estimate the optimal policy. One particular algorithm is called *value iteration*. This is an iterative algorithm which starts with any randomly initialized value function V_0 . With each iteration k , the algorithm produces an approximation V_k , until it converges to the optimal value function V^* .

The value function is updated according to the following update rule:

$$V_{k+1}(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')], \quad \text{for all } s \in S,$$

where $T(s, a, s')$ is the probability to transition to the state s' given state s and action a , and $R(s, a, s')$ is the immediate reward for transition to the state s' given state s and action a .

Here index k can be viewed as the number of future rewards ahead that we are taking into account. For example, $k = 0$ means that we are considering 0 time steps ahead, meaning that we are not considering even immediate reward.

Policy Iteration.

One of the disadvantages of the value iteration algorithm is due to the stopping criteria of the algorithm. In order to find whether we should stop the algorithm we need to observe the values for iterations k and $k + 1$. We can stop the algorithm only when we have a significant evidence that the values for all states are no longer improving. This gives rise to the situation when the optimal action for each state is already identifiable but there is still some difference between V_k and V_{k+1} , for example due to small learning rate. Thus the algorithm may spend unnecessary amount of time just to meet formal criteria of convergence.

The algorithm called *policy iteration* (Bellman, 1956) is able to address this particular problem. The first step of the policy iteration algorithm is called the *policy evaluation* (Bellman, 1956). During the policy evaluation step, the algorithm is given any intermediate policy π_i . This policy is then evaluated using the following update rule:

$$V_{k+1}^{\pi_i}(s) = \sum_{s'} T(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s')], \quad \text{for all } s \in S,$$

where $T(s, \pi_i(s), s')$ is the probability of transition to the state s' given state s and action $\pi_i(s)$, and $R(s, \pi_i(s), s')$ is the immediate reward for transition to the state s' given state s and action $\pi_i(s)$.

The policy evaluation step has a lot of similarities with the value iteration algorithm. Only in the policy evaluation we are not looking for a maximum over all possible actions and we are using a fixed policy π_i .

Once the policy evaluation has converged, we proceed to the *policy improvement* (Bellman, 1956) step. We update the intermediate policy π_i according to the following rule:

$$\pi_{i+1}(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k^{\pi_i}(s')], \quad \text{for all } s \in S,$$

where $T(s, a, s')$ is the probability to transition to the state s' given state s and action a , and $R(s, a, s')$ is the immediate reward for transition to the state s' given state s and action a .

The stopping criteria for the policy iteration is reached when we no longer able to improve intermediate policy π_i and thus we declare π_i to be the optimal policy π^* .

The advantage of the policy iteration over the value iteration lies within the policy evaluation update rule. This update rule doesn't have the maximum over all action operator (\max_a), since the policy under evaluation is fixed; This gives a significant speed of convergence advantage.

1.1.5 Q-Learning.

Policy iteration doesn't directly translates to the realms of the reinforcement learning as it requires the knowledge of a model of an environment (knowledge of reward and transition functions). Look at the update rule for the policy evaluation:

$$V_{k+1}^{\pi_i}(s) = \sum_{s'} T(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s')].$$

And the update rule for the policy improvement:

$$\pi_{i+1}(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k^{\pi_i}(s')].$$

They both require the knowledge of the functions T and R , where T is transition function and R is a reward function.

The reinforcement learning equivalent of policy iteration is called *Q-Learning* (Watkins, 1992). Q-Learning, like DP value iteration, makes sure that any randomly initialized Q function converges to optimal action-value function Q^* . See Algorithm 1 for details.

Q-Learning is a temporal-difference learning algorithm (Sutton and Barto, 1990). The value iteration algorithm uses an expected value operator which requires knowledge of the state-transition function. Temporal-difference learning algorithms work around this by modifying update rule to work with raw experience. The update rule is slightly modified but essentially it does the same job only without the expected value operator:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right].$$

Q-Learning requires several criteria to be met:

- Underlying problem must be a Markov Decision Process.
- Learning rate α should decrease to zero.
- Every state-action pair should be tried infinitely many times.

However, in practice the algorithm converges without some of this strict criteria. For example, in many cases the algorithm is able to converge with fixed α parameter.

1.1.6 Off-policy and On-policy algorithms.

One can run a series of episodes starting from the same starting state and record the return in each of the episodes. The average of those returns will differ from the value of

Algorithm 1 Q-Learning

Input: Initialize $Q(s, a)$ randomly
repeat
 Initialize starting state s
 repeat
 Choose a from s using policy derived from Q
 Take action a , observe r, s'
 $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
 $s \leftarrow s'$
 until s is terminal
until no more episodes

the starting state predicted by learnt optimal Q-function. That is there will be a difference between the learned optimal Q-function and the actual results while the agent is learning. This behaviour is due to:

- Q-learning learns the value of optimal sequence of actions given any starting state.
- The algorithm will introduce a randomness into action-selection process to insure that every state-action pair is selected (required by proof of convergence). Thus the actual sequence of actions taken by the agent will almost never be optimal.

Off-policy algorithms, such as Q-learning, do not record the actual value of each state-action pair; They miss the fact that during each action-selection phase there might be a random action with potential negative outcome. As a result the agent might easily identify the greediest action among available actions but it will miss the fact that taking that greedy action might potentially result in big penalties.

On the other hand, on-policy methods do take into account the randomness associated with the action selection mechanism of the policy being followed. Thus the values $Q_\pi(s, a)$, for each state action pair, reflect the average of all sequences of actions that are possible under the policy π .

The Q-learning is ideal for applications where an identification of optimal policy is the only goal. That is when there is no harm done if the agent has performed a particularly bad sequence of actions. An example is escaping from a maze - the goal is to find an

optimal escape route. It doesn't matter how many times the agent was circling around the maze while exploring. On the other hand, if a Mars rover is exploring the surface of the planet, it becomes very expensive and potentially dangerous for the equipment.

1.1.7 S.A.R.S.A.

Sarsa is the on-policy temporal-difference algorithm similar to Q-learning. Both algorithms are very similar; The only difference is the update rule employed in each of the algorithms. The Q-learning has the following update rule:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)].$$

Here the construct $\max_{a'} Q(s', a')$ reflects the fact that only the greediest path from successor state s' is taken into account while updating the state-action pair $Q(s, a)$.

The Sarsa algorithm has the following update rule:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)].$$

Sarsa doesn't try to maximize the value with which to update $Q(s, a)$. Instead, $Q(s', a')$ is the return from successor state after taking action a' , which was selected according to the policy being followed by the agent.

1.2 Multi-objective Optimization.

Multi-objective reinforcement learning methods are built upon methods from general multi-objective optimization. Therefore this section will briefly describe the main concepts and techniques that were borrowed from general multi-objective optimization.

1.2.1 Multi-objective Problem.

Let A be a set of all possible actions. Given n objective functions $f_1(), \dots, f_n()$, we map each action a from the set A to a point $(f_1(a), \dots, f_n(a))$ in the n -dimensional objective space.

Multi-objective problem is the problem of finding actions a from the set A such that we are most satisfied with the reward vectors $(f_1(a), \dots, f_n(a))$.

1.2.2 Tradeoffs.

Multi-objective optimization aims to solve difficult real life problems and many of these problems involve objectives that conflict with each other. Introduction of these conflicting objectives means that in many cases we will not have a single best option. Or in other words we cannot improve all objective functions at the same time. Rather we will need to compromise between two or more objectives, gaining more value on one objective at the expense of the other(s).

1.2.3 Dominance.

Given two solutions a' and a'' we can map them into two points in the n -dimensional objective space:

$$x' = (x'_1, \dots, x'_n) \quad \text{and} \quad x'' = (x''_1, \dots, x''_n)$$

where

$$x'_i = f_i(a') \quad \text{and} \quad x''_i = f_i(a'') \quad \text{for} \quad i = 1, \dots, n$$

The point x' is said to Pareto dominate (Pareto, 1896) point x'' when

$$x'_i \geq x''_i \quad \text{for all } i,$$

and

$$x'_i > x''_i \quad \text{for at least one } i.^1$$

Dominated points by definition will imply the existence of solutions that achieve more on each of the objectives and because our goal is to maximize the objective function we may omit dominated solutions. On the other hand, we may find two solutions which outperform each other on one of the objectives. In this case we really cannot say which solution is better, since each of them excels on one of the objectives. We call them non-comparable solutions. The set of all non-comparable solutions constitutes the Pareto front (Pareto, 1896).

1.2.4 The Pareto Front.

Every action a in the actions space A is getting mapped to a point $(f_1(a), \dots, f_n(a))$. The image of the set A in the n -dimensional objective space is called *the range set* R . The subset of R which consists of non-dominated solutions is called *the Pareto front*, as shown in Figure 1.2.

Now we can reformulate multi-objective problem as the problem of locating any of the points belonging to the Pareto Front. Which point should it be? All of the points are not comparable to each other. If the tradeoff is a matter of a personal preferences then a personal opinion is required to make a decision.

1.2.5 Selecting a solution in the Pareto Front.

To reach a point in the Pareto front we need a scalar index which will reflect the preference between any two points. We can introduce a function v defined on the *range set* R with the following property:

$$v(x_1, \dots, x_n) \geq v(y_1, \dots, y_n)$$

¹For consistency with reinforcement learning discussion it will be assumed that the agent's objective is to maximize the objective functions, whereas in most multi-objective optimization literature the goal is to minimize the objectives. In fact, the two cases are fundamentally equivalent.




Figure 1.2: The Pareto front is represented with black points. Each black point is a non-dominated solution. Grey points represent dominated solutions. Each gray point is dominated by at least one black point.

if and only if the point (x_1, \dots, x_n) is more preferable or equal to the point (y_1, \dots, y_n) .

The following two sections will briefly describe two techniques to introduce that function v . These techniques are used in a large portion of multi-objective reinforcement learning algorithms and thus deserve a mention.

Linear-Weighted Averages.

Assume as before that the existence of a solution set A and we also have n objective functions $f_1(), \dots, f_n()$. Each solution $a \in A$ is mapped to a point $(f_1(a), \dots, f_n(a))$. Let

$$\mathbf{w} = (w_1, \dots, w_n)$$

be a weight vector. Where

$$w_i > 0, \text{ for all } i$$

and

$$\sum_{i=1}^n w_i = 1.$$

Let us pose an additional optimization problem: choose $a \in A$ that will maximize

$$\sum_{i=1}^n w_i f_i(a).$$

Assume solution a^* maximizes this given optimization problem. Let $\mathbf{x} = (f_1(a^*), \dots, f_n(a^*))$, in the n -dimensional objective space, with solution a^* .

The claim is that this point \mathbf{x} must belong to the Pareto optimal set. If this point was not a member of the Pareto front there must another point \mathbf{y} which dominates point \mathbf{x} . But it cannot be since it would imply

$$\sum_{i=1}^n w_i y_i > \sum_{i=1}^n w_i x_i$$

And we already know that point \mathbf{x} is the maximum of the given optimization problem.

The more weight we apply to the specified objective the more favourable to the specified objective will be the identified point of the Pareto front. This allows us to perceive weights as a form of preference. However, setting the preferences a priori means human interaction and sometimes it is not so clear how to set these preferences.

Another limitation of these approach is the inability to locate concave points of the Pareto front. For example if the Pareto front has a concave shape then this approach will only locate extreme points.

Lexicographical Ordering.

Another widely used approach is to order objective functions $f_1(), \dots, f_n()$ according to their importance.

Given two solutions a^1 and a^2 we say that a^1 is better than a^2 when

$$f_i(a^1) = f_i(a^2), \quad \text{for } i = 1, \dots, n-1$$

and

$$f_n(a^1) > f_n(a^2).$$

This approach is intuitively easy to understand and in fact in many problems it is possible to clearly differentiate which objective is more important.

1.3 Multi-objective Reinforcement Learning Research.

Reinforcement learning initially was applied to find optimal policies in single objective problems. However, in recent years there was a significant increase in the interest towards multi-objective reinforcement learning. This can be explained by the fact that sometimes it is just impossible to express a sophisticated problem in terms of just single objective, and naturally this gave rise to the multi-objective reinforcement learning. For further motivation refer to Roijers et al. (2013).

This section will review the main work done on temporal-difference multi-objective reinforcement learning. It is only a review of the most seminal works, for more details please refer to Roijers et al. (2013). First subsection will review temporal-difference methods based on the linear scalarization; Second subsection will consider non-linear methods. This separation is caused partially because linear methods are proven to converge while non-linear may not.

1.3.1 Linear Temporal Difference Learning.

Before linear scalarization was used in reinforcement learning, Feinberg and Shwartz (1993) wrote an extensive paper on using linear scalarization in constrained markov decision models. The authors proved the existence of optimal policies in such models and also proved that linear scalarization is able to identify members of the Pareto front. Even though this paper didn't mention reinforcement learning, it provided the necessary theoretical foundations for other researchers to try out linear scalarization in multi-objective reinforcement learning.

Karlsson (1997), in his PhD thesis, first tried to combine Q-learning algorithm and linear scalarization. The author argues that when dealing with multi-criteria problem, it is much better to break a complex state space into a number of smaller state spaces, so called modules. For each module we maintain a separate Q-function. It helps to avoid relearning when dealing with a complex task consisting of a number of independent sub-tasks.

Having multiple separate Q-functions, Karlsson faced a problem of how to administer the action selection process. He used the so called "greatest mass method". To determine a value of some action under this method, one needs to add this action's Q-values from all modules. An action with the greatest overall sum is pronounced the greedy action. Essentially, this method is just a linear scalarization with all weights equal to one.

The author provided details of his test environment, which he called the simple driving domain (SDD). The SDD is a grid world with roads crossing it vertically and horizontally. The world is populated with different objects: street signs, traffic lights and obstacles. An agent controls a car. The agent has four goals: stay on a road, avoid obstacles, negotiate streetlights, follow street signs. All tasks are independent of each other. The author specifically created independent sub-tasks to test the modular approach.

The main point of Karlsson's research was to investigate how agents behave when the complex state space is decomposed into smaller sub-spaces. For this purpose he also created one agent which maintained a complex structure of the state space (monolithic approach). His results show that in many scenarios the modular approach performs as well as monolithic approach, while having a significant advantage in terms of speed of convergence.

Castelleti et al. (2002) further explored linear scalarization. They were interested in applying reinforcement learning methods to a water reservoir management. As a case study they chose Lake Como, which located in Northern Italy. This lake provides water supply to agricultural districts and hydroelectrical power plants.

It is possible to model this problem as an MDP. State information includes: current water level and information about natural water income. Actions govern how much water to release from the reservoir. In this paper the authors used two objectives: provide water and avoid flooding.

Unlike the work of (1997), the researchers didn't break the complex state space into smaller ones. However essentially their approach is the same - to identify greedy actions, one combines Q-values of each objective into a weighted sum. Only in this case, weighted sum is no longer uniform and weights could be arbitrary (as long as they sum to 1). They called their algorithm Qlp (Q-Learning planning). The authors approximated the Pareto front by having multiple independent tests - each test for one weight set.

In the experimental results section the authors compared their proposed approach with the results obtained with the help of stochastic dynamic programming approach (SDP); The QLP consistently outperformed the SDP.

It is important to notice that it was probably the earliest work where non-uniform weights were used.

Natarajan and Tadepalli (2005) used linear scalarization to sequentially learn all optimal policies. All the papers discussed so far organized trials in similar manner:

1. Set preferences for each objective.
2. Initialize value function.
3. Learn optimal policy for a the specified set of weights.

All these steps need to be performed when preferences are changed. The authors noticed that the learnt value function for previous set of weights can be used to speed-up the learning process. If sets of weights are significantly different then one will not observe noticeable improvements. On the other hand, if the preferences correspond to nearby points in the objective space then the learning process will take full advantage of the value function learnt before. The authors propose mechanism which allows to determine which of the previously learnt policies should be used when the preferences are changed.

Two experimental domains are included in the paper: modified Buridans ass problem and network routing problem. The modified Buridans ass is the modification of the original problem (Gábor et al., 1998). The modified version has an increased state-space - 3 by 3 grid world. A donkey is located in the middle of the grid world and two plates of food are located in top-left corner and bottom-right corner of the world respectively. The donkey must reach either of the plates with food to stay alive but every time he reaches for one of the plates the one might be stolen. Overall there are three objectives: amount of food eaten, number of plates stolen and a number of moves made by the donkey. The authors demonstrate experimental results, which show a significant reduction of time-steps needed to learn optimal policies when prior value functions.

The network routing domain represents typical computer network with a number of nodes connected through a specified network topology. Three objectives are present: end to end delay, number of packets lost and power consumed by a node. Each state is represented by a current node and destination node. Actions are represented by neighbors of the current node through which the packet might be sent. Overall goal is to find an optimal routing policy according to specified preferences. The experimental results again demonstrate that learning from scratch for each set of weights takes a significantly longer amount of time than learning from prior policies.

1.3.2 Non-linear Temporal Difference Learning.

In many applications there is a natural ordering of the objectives; For example, consider a cleaning robot with two objectives: maintain a battery charged at all times, and clean an office. In this scenario there is natural ordering - the battery of the robot should never deplete (cleaning is not possible if the robot's battery is depleted), thus maintaining the battery level is considered more important. Non-linear temporal-difference methods can accurately reflect this ordering of the objectives and thus are more suitable for this domain.

More importantly, it has been shown (Yaman et al., 2010) that people often employ lexicographic ordering of the objectives while making complex decisions.

One of the most earliest and most widely cited publications in MORL is Gabor, Kalmar and Szepesvari (1998). The publication first develops a solution for multi-criteria dynamic programming. After that the solution is modified to be extended to multi-criteria reinforcement learning.

The authors aimed at developing a mechanism to order policies. For this purposes, they used a minimum threshold on one of the objectives. This allowed authors to use only one objective as a maximization criteria. This paper builds upon lexicographical ordering of objective functions, which was described in the the multi-objective optimization section (Section 1.2.5).

After showing that this approach converges for abstract multi-criteria dynamic programming problem, the authors extend the technique for multi-criteria reinforcement learning.

This publication goes to a great extent describing mathematical proof of the proposed algorithm. Apart from mathematical aspects, the authors also included experimental results. Using a game of tic-tac-toe, the authors were able to show that their algorithm was able to identify optimal policy. They used two objectives: to win a game and to do it as fast as possible.

Van Moffaert et al. (2012) proposed to use Chebyshev scalarization function. They used Chebyshev distance in the process of evaluating a set of actions. To successfully use this distance the authors introduced a so called utopian point z^* in the objective space. After applying scalarization function to any state-action combination (s, a) we obtain a scalarized Q-value (SQ):

$$SQ(s, a) = \max_{o=1\dots m} w_o \cdot |Q(s, a, o) - z_o^*|$$

, where m is the number of objectives. Given any state s and all available actions a , we define a greedy action to be a solution to the following optimization problem:

$$\min_a SQ(s, a).$$

The authors tested their approach on two well studied multi-objective reinforcement learning problems: the Deep Sea Treasure and the multi-objective Mountain Car. They were able to demonstrate that on both problems their non-linear scalarization approach was able to locate more optimal policies than usual linear scalarization.

Van Moffaert et al. (2013) proposed to use hypervolume indicator to guide action selection process. Their action-selection strategy (algorithm 2) drastically differs from both linear and non-linear scalarization methods.

The agent maintains a list l of points in the objective space. The list l is empty at the beginning of an episode. On every step of the episode we map selected state-action to the objective space, according to q-value of the state-action pair. The resulting point is added to the list l .

Given any state s and some available action a , the proposed strategy maps this action to a point in the objective space. After that we add this point to the list l and calculate the hypervolume. The action is called greedy when it has the biggest contribution to the hypervolume.

The list l is empty at the beginning of an episode. On every step of the episode we map selected state-action to the objective space, according to q-value of the state-action pair. The resulting point is added to the list l .

This approach utilizes search in the objective space of a problem. That is the main difference from the methods based on the scalarization, where we take q-values for each objective and transform it to a single value. However, this algorithm also uses Q-learning as underlying learning mechanism.

Algorithm 2 Hypervolume based action-selection

Input: a state s , list l of Q-values
 $volumes \leftarrow \{\}$
for each action a available at given state s **do**
 create m-component vector \mathbf{r}
 for $i = 0$ to m **do**
 $r_i \leftarrow Q(s, a)$ for objective i
 end for
 $hv \leftarrow \text{getHypverVolume}(l + \mathbf{r})$
 add hv to $volumes$
end for
return $\text{argmax}_a volumes$

The authors used the same set of test problems as in their work with Chebyshev scalarization function. On the Deep Sea Treasure problem, the hypervolume-based algorithm slightly outperformed one based on Chebyshev scalarization function, and significantly overpowered linear scalarization. On the multi-objective Mountain Car problem, the new algorithm performed comparably to Chebyshev scalarization function, and significantly outperformed one based on linear scalarization.

We know from Roijers et al. (2009) that TD methods using non-linear scalarization may fail to converge due to the non-additive structure of the returns. While the work discussed in this section reports successful results based on non-linear temporal-difference, it may be that this is due to the specific characteristics of the small number of test problems on which performance has been evaluated. A need exists for more comprehensive evaluation of such methods.

1.3.3 Simultaneous Learning of More Than One Policy.

The algorithms reviewed in previous two sections belong to the class of single-policy methods, in that each run of the algorithm finds a single Pareto optimal policy (Vamplew et al., 2011). To find an approximated Pareto front it is required to run the algorithm multiple times with different parameter setting (one for each Pareto optimal policy) and then

manually combine them. In contrast, multi-policy algorithms aim to identify an approximated Pareto front in single run. This section will examine most important research in this direction.

Barrett et al. (2008) propose an idea which allows to learn all optimal policies at the same time. They exploit the fact that Pareto front is essentially a convex hull object, under the assumption that linear scalarization is used to identify optimal points. This allows us to use convex hull operations in multi-objective reinforcement learning realms.

Single objective problems are naturally suitable for the Q-Learning update rule

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right].$$

Here we are backing a Q-value for previous state-action pair with immediate reward and maximum expected reward from successor state. In single-objective problems, it is always easy to find which sequence of actions will gain more rewards. So it is always clear which value to use in this update rule. In contrast, multi-objective problems lack this clarity of choice. Now, we may have a multiple number of incomparable sequences of actions. The authors propose that we should use all optimal sequences of actions from a successor state and use all of them to back previous state-action pair. We no longer keep a single value for any state-action pair, rather we keep a set of values (convex hull of optimal Q-values). Thus we keep track of all possibilities for each state-action pair. At any time we can apply preferences (represented as a set of weights) to identify optimal policy.

The paper contains proof of correctness of this approach. The authors also included one test problem called Resource Gathering. This problem is a two dimensional grid world, an agent can locate two types of resources in this grid world. The agent receives a reward for each item brought home. However there are also two enemy positions where the agent can be attacked. If attacked, the agent receives a penalty and is moved to the home state.

This paper proposes an idea of learning optimal policies without ever specifying our preferences. This is a completely opposite approach to most of the research done. Usually, we first set preferences and only after that we learn policy optimal to those preferences.

Lizotte et al. (2010) also propose a mechanism to learn all optimal policies. They work with data from randomized medical treatments. These are the trials set up to find out how different patients react to different sequences of trials. The set of available actions consists of all possible treatments and set of all states consists of patient's health observations. They concentrate on two-objectives scenario, however, they do propose suggestions on how to extend their approach to more than two objectives. The first objective is the severity of the symptoms and the second objective measures the severity of the side-effects.

They use model-based approach; Originally state-transition probabilities and reward function are not known for apparent reasons but it is possible to analyze data from the trials and make approximations for both state-transition probabilities and reward function. Each episode:

$$s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_T, a_T, r_T$$

is a history of prescribed treatments for one particular patient. At each time step t , s_t is the health observations and a_t is the prescribed treatment at this time step. These trajectories are used to make a model of the MDP.

After the the state-transition probabilities and the reward function are approximated, the authors use bottom-up dynamic programming approach. First they calculate the Q-function for the last time step T :

$$Q_T(s, a, \delta) = r_T(s, a, \delta),$$

where $r(s, a, \delta)$ is

$$r(s, a, \delta) = (1 - \delta) \cdot r_0(s, a) + \delta \cdot r_1(s, a),$$

where $\delta \in [0, 1]$, $r_0(s, a)$ returns immediate reward for the first objective and $r_1(s, a)$ returns immediate reward for the second objective.

After calculation of Q_T it is possible to calculate the Q-function for a time step t that precedes time step T :

$$Q_t(s, a, \delta) = r_t(s, a, \delta) + E_{s'|s, a} [V_{t+1}(s', \delta)],$$

where $V(s, \delta) = \max_a Q(s, a, \delta)$.

Modifying delta essentially changes preferences between the objectives and allows us to target different individual points of the Pareto front. The authors noticed that it is possible to choose a few values of δ and calculate their respective representations in the objective space. After that it is possible to use linear spline interpolation to approximate all other points.

The authors compare their approach to the Barrett et al. (2008) and show that asymptotically their approach is better in terms of time and space complexity.

Chapter 2

Empirical Evaluation.

2.1 Introduction.

To conduct an empirical study one needs means to compare different algorithms. For example for classification algorithms there are repositories of data sets, with each data set describing a particular problem. These data sets allow researches to identify weaknesses or advantages of their algorithms. They also allow them to make cross-comparison between different algorithms.

Reinforcement learning in general and multi-objective reinforcement learning in particular are very different from research areas like classification or clustering. Reinforcement learning is a multi-step decision making process with a number of actions available at each state at every time step. Add to this a factor of randomness associated with state transition and reward functions. Having all these factors substantially complicates the creation of a test problem. In reinforcement learning a test problem is not a static data set but rather a software which reacts dynamically to the actions of an agent.

The absence of central repository of test problems lead to the fact that every published paper is disconnected from all other papers in terms of test problems. All papers described in literature review present a brand new test environments. Moreover the majority of the papers use only small amount of test problems, usually two. This is in contrast with other

areas of machine learning where each published paper is connected with other papers by using the same test data sets and the number of tested data sets is usually high.

Last section of the first chapter reviewed some of the most important research in temporal difference multi-objective reinforcement learning. All reviewed papers were mainly concerned with proving the convergence of their algorithms. In each paper authors presented a small number of their own test problems and their own methodology. It is enough in terms of convergence proof but clearly not enough to provide the full guidance about the algorithm. The following information may be included in such a guidance:

- Information about the structure of a reward function. Whether the reward is intrinsic or extrinsic may have a profound effect on the performance of some algorithms.
- Each test problem must be built around particular property, which can be observed in realistic domains, such as the shape of the Pareto front.
- Each algorithm should be subjected to a number of different test problems, to identify a realistic domain of application.

One can also notice the absence of common methodology across all reviewed papers. Thus one cannot easily compare the results of an algorithm presented in one paper with the results of different algorithm. Each methodology essentially incorporates a unique idea of how to show whether an algorithm is effective or not. Some of those ideas may drastically differ. Overall it substantially complicates the process of comparing the results presented in different papers.

The arguments above can explain the absence of empirical results so far. No single attempt was made to perform empirical evaluation of the proposed algorithms. Some of the papers provided empirical evaluations but only to prior versions of the same algorithms.

One of the difficulties associated with an empirical evaluation of existing algorithms is that different methodologies are exploited by different researchers. One major example can clarify the point. In multi-objective optimization the quality of the algorithm is evaluated by how close the approximated Pareto front is to the true Pareto front of the

problem. Currently, there is a vast number of proposed performance metrics which allow to assess how close the solution is to the true one. Clearly, it is impossible to choose from two algorithms when you know that the results provided in a published paper were obtained with different set of performance metrics in each paper.

2.2 The Details of the Empirical Evaluation Study.

All of the works reviewed in the literature review chapter have the same limitations:

- Limited number of test problems.
- No mention about the inability of the linear scalarization based algorithms to find concave points of the Pareto front. This limitation was quite known in traditional multi-objective optimization but until recently (Vamplew et al., 2008) was never specifically mentioned in reinforcement learning theory. As a result later research work with uniform weights (Ferreira et al., 2012, Aissani, Beldjilali, and Trentesaux (2008), Shabani (2009)) also didn't mention possible limitations of this approach.
- Each author or team of authors used their own evaluating methodology which complicates cross-comparison.

This document will use the methodology previously proposed by Vamplew et al. (2011). This paper outlined very important aspects that should be introduced:

- MORL algorithms should be tested on a wider range of test problems, and the properties of these problems should be understood to aid in interpreting the variations in the performance of algorithms.
- Standard benchmark problems and standard implementations of these problems should be established so as to facilitate comparison of the performance of different algorithms.

- Standard approaches to experimental methodology and reporting of results should be adopted, again to aid in the comparison of algorithms between papers. In particular numeric measures of performance will prove more useful for this purpose than the graphical reporting of results.

2.2.1 Performance metrics.

Every multi-objective problem implies existence of a set of Pareto optimal solutions. The task of the learning algorithm is to find an approximation to that front. The quality of the approximation defines the quality of the algorithm. Although every task has a true set of Pareto optimal solutions, machine learning algorithms can usually only approximate this set. And every algorithm does it differently, capturing specific points of the true front. Thus performance metrics should be introduced to facilitate in assessing the quality of the found set of Pareto optimal solutions. Multi-objective reinforcement learning is similar to general multi-objective optimization in terms of a set of Pareto optimal solutions thus previous results in multi-objective optimization could easily be used in MORL.

Each approximation to a set of Pareto optimal solutions captures only particular points of the true set. How these points are spread between each other? How many points are captured? And to which extent does the approximation covers the true set? All these important features could be encapsulated into a metric which allows us to compare two different learning algorithms. Early work in multi-objective optimization used scalar representation of these features. But each scalar representation captures only specific detail of the front such as diversity, cardinality or accuracy.

Later work in multi-objective optimization used different approach. Instead of multiple number of scalar metrics, a composite metrics compatible with the notion of Pareto dominance were proposed. Berry (2008)) showed a preference of such composite metrics over a number of scalar ones. A number of such approaches were used in multi-objective optimization.

Hypervolume metric.

The hypervolume metric (Zitzler and Thiele, 1999) is a well-recommended composite metric which is used to evaluate the quality of an approximate Pareto front found by a learning algorithm. It measures the volume of the objective space region which was created by the points in the found Pareto front approximation. One additional point is also added before the volume is measured. This point is dominated by all points in the Pareto front approximation and serves as a reference point. An example is shown in figure 2.1. The reference point can be used as a tuning mechanism and allows greater flexibility.

Given two learning algorithms, we can run them on the same test problem. After that we get a Pareto front approximation from each algorithm and calculate the hypervolume for each approximation. The approximation which produced a higher value of the hypervolume is a better approximation to the true Pareto front. This is proven to be true both for convergence and diversity.




Figure 2.1: The figure shows how points from a Pareto front approximation and a reference point together form the region of objective space. The black dots are Pareto optimal solutions and the red dot is the reference point.

One important remark - the reference point must be consistent when evaluating approximations obtained by different learning algorithms.

2.2.2 Experiment Structure.

Single-Policy Algorithms.

Single-policy algorithms (Castelleti et al., 2002 or Gabor, Kalmar and Szepesvari, 1998) are initialized with a set of preferences. After that the algorithm executes a number of learning episodes and hopefully it will learn the Q-function of the optimal policy for the specified preferences.

To illustrate the work of single policy algorithms consider the following example; A simple MDP was constructed, it consists only of three time steps. On every time step two actions are available - left and right. State transition is deterministic i.e. each action always leads to the same successor state. For more details refer to Figure 2.2.




Figure 2.2: Blue circles represent states; Black arrows represent the actions. In every state left and right actions are available. State transition is deterministic and an action always leads to the same next state. Next to every arrow there is a vector of immediate rewards for choosing this action. The overall return starting from starting state is shown at the bottom of each branch.

One can identify four different policies for this example MDP. Essentially, every branch of this tree is a unique policy with it's associated return.

It is possible to look at this problem from a different perspective - from the objective space. Figure 2.3 shows two dimensional objective space for the example MDP problem.

Each blue dot corresponds to a unique policy and its coordinates are exactly this policy's return. Overall there are four points in the objective space, one for each policy. One can notice that these four points together constitute the Pareto front for this problem. This is because the example MDP problem doesn't have dominated policies.




Figure 2.3: Each blue dot is a unique policy from the example MDP problem. **NOTE:** Because the example MDP does not have dominated policies all four available policies together constitute the Pareto front for this problem.

Consider any single-policy algorithm, for example one discussed in Castelletti et al. (2002). This algorithm requires a set of weights to be specified before the algorithm can start learning; Let $\mathbf{w} = \{0.4, 0.6\}$ be this set of weights. Essentially, this algorithm is the multi-objective reinforcement learning adaptation of the linear-weighted average algorithm from general multi-objective optimization. Linear-weighted average algorithm will find which of the available solutions (policies) will maximize the dot product of the return vector and the weight vector. In case of the example MDP, the policy with the return of $[60, 60]$ (see Figure 2.4) will maximize the weighted average. The single-policy algorithm will learn the Q-function associated with this policy.




Figure 2.4: The blue dot which represents the policy with return $[60,60]$ is marked with red circle around. The single policy algorithm mentioned above will learn the Q-function for the policy associated with the marked blue dot.

Even though the algorithm was able to identify one element of the Pareto front, it is still preferable to find the whole approximation of the front. The only way to achieve this with single policy algorithms is to reset the preferences every time the algorithm tries to identify another member of the Pareto front. This requires a special experiment design:

1. Initialize an array of preferences, where each preference is an array itself (for example a set of weights).
2. For each item in the array of preferences run a single-policy algorithm for a specified number of episodes. After the algorithm has converged, extract the greediest Q-value of a starting state and record it. **NOTE:** Notice that a priori an experimenter cannot identify which member of the Pareto front will be targeted by a certain set of preferences. One can only wait until the algorithm has learned the Q-function of a policy which is associated with the set of preferences. After that the greediest Q-value of the starting state will identify which point in the Pareto front was identified.
3. Combine all recorded points into an approximation of the Pareto front.

This experiment setting will allow to create an effective approximation of the Pareto front of an underlying problem. It is vital to have this approximation if one is to compare the performances of different algorithms.

Multi-Policy Algorithms.

Single-policy algorithms are initialized with a set of parameters, usually its a set of weights - one for each objective. Eventually, this set of parameters gets mapped into one of the policies and the algorithm will learn the Q-function for this policy. As was mentioned before, an experimenter needs to prepare a number of sets of preferences (weights), and hopefully those sets will approximate the Pareto front and restart the experiment for every set of preferences.

This experiment setting is not very efficient - not only it takes more computer time but also significantly complicates an experiment design. Instead of sequential working with each set of parameters one can work with all sets of parameters at the same time. For example, an algorithm might be given the following set - $\{ \{0.7, 0.3\}, \{0.3, 0.7\} \}$. Here the algorithm works with two sets of weights. Each set of weights represents the relative importance of each objective to an experimenter. Ideally, both sets will get mapped into different policies, although they both might get mapped into the same policy. The algorithm then should learn the Q-functions for those mapped policies.

In this setting, one of the mapped policies is used as an active policy; It's Q-function is used during action-selection period. Other mapped policies are used only for learning purposes.

Simultaneous learning of multiple policies is possible because state-transition dynamics provides enough information for more than one policy to be successfully updated. As an example consider state-transition depicted on Figure 2.5. Let $\vec{w}_1 = \{0.7, 0.3\}$ and $\vec{w}_2 = \{0.3, 0.7\}$ be two weight vectors. Each vector represents relative importance of objectives. Weight vectors \vec{w}_1 and \vec{w}_2 will get mapped to policies $\pi_{\vec{w}_1}$ and $\pi_{\vec{w}_2}$ respectively. Let $\pi_{\vec{w}_1}$ and $\pi_{\vec{w}_2}$ be deterministic policies.

Further let $\pi_{\vec{w}_1}$ be an active policy. At time step t action-selection was made according to active policy $\pi_{\vec{w}_1}$ and let $a_t = \text{left}$. Pair (s_t, a_t) makes the environment to transition to state s_{t+1} and the immediate reward $[20,20]$ is received by the learning algorithm.




Figure 2.5: The blue dot which represents the policy with return $[60,60]$ is marked with red circle around. The single policy algorithm mentioned above will learn the Q-function for the policy associated with the marked blue dot.

At time step $t + 1$ it is time to update $Q(s_t, a_t)$ for both $\pi_{\vec{w}_1}$ and $\pi_{\vec{w}_2}$. To update $Q(s_t, a_t)$ the algorithm will need to identify greedy action for the state s_{t+1} and use Q-value of that greedy action in combination with the immediate reward. For each policy the greedy action might be different but the immediate reward will be the same. This allows to update Q-function for the policy $\pi_{\vec{w}_2}$ even though the state-transition was generated by the policy $\pi_{\vec{w}_1}$.

Table 2.1 shows the Q-functions for the state s_{t+1} for both policies $\pi_{\vec{w}_1}$ and $\pi_{\vec{w}_2}$. These two Q-functions are identical because both actions lead to the terminal state and only immediate reward is used during the update. For each policy the greedy action is the action which will maximize the weighted sum:

$$\sum_{i=1}^n w_i f_i(a).$$

For the policy $\pi_{\vec{w}^1}$ the weight vector is $\vec{w}^1 = \{0.7, 0.3\}$ and thus the greedy action will be *left* because

$$(60 \times 0.7) + (40 \times 0.3) > (40 \times 0.7) + (60 \times 0.3).$$

On the other hand for the policy $\pi_{\vec{w}^2}$ the weight vector is $\vec{w}^2 = \{0.3, 0.7\}$ and the greedy action will be *right* because

$$(60 \times 0.3) + (40 \times 0.7) < (40 \times 0.3) + (60 \times 0.7).$$

Thus the update rule for $Q^{\pi_{\vec{w}^1}}(s_t, a_t)$ will take the following form:

$$Q^{\pi_{\vec{w}^1}}(s_t, a_t) \leftarrow Q^{\pi_{\vec{w}^1}}(s_t, a_t) + \alpha \left[r_{t+1} + \gamma Q^{\pi_{\vec{w}^1}}(s_{t+1}, \text{left}) - Q^{\pi_{\vec{w}^1}}(s_t, a_t) \right].$$

And the update rule for $Q^{\pi_{\vec{w}^2}}(s_t, a_t)$ will take the following form:

$$Q^{\pi_{\vec{w}^2}}(s_t, a_t) \leftarrow Q^{\pi_{\vec{w}^2}}(s_t, a_t) + \alpha \left[r_{t+1} + \gamma Q^{\pi_{\vec{w}^2}}(s_{t+1}, \text{right}) - Q^{\pi_{\vec{w}^2}}(s_t, a_t) \right].$$

Table 2.1: **Left-hand side table** is the Q-function for the state s_{t+1} for the policy $\pi_{\vec{w}^1}$. **Right-hand side table** is the Q-function for the state s_{t+1} for the policy $\pi_{\vec{w}^2}$.

state \ action	left	right	state \ action	left	right
s_{t+1}	[60,40]	[40,60]	s_{t+1}	[60,40]	[40,60]

2.2.3 Benchmarking Software.

As was mentioned before a test problem in reinforcement learning is a dynamic software, which actively responds and changes according to the actions from an agent. The agent is also of dynamic nature and in the similar manner dynamically reacts to the changes in the test problem. This due to the inherent dynamic nature of the the agent-environment cycle. For example every test problem should have the means(function or method) to

accept an action from the agent. As well as means (again function or a method) to pass information about state-transition and reward functions.

Even though test problems may have nothing in common in terms of realistic domains they represent, they still share those means of interaction with an agent or learning algorithm. Moreover in every problem the information that is passed between an agent and an environment is of very similar nature: the agent informs the environment about his preferred action, the environment notifies the agent about the consequences of that particular action.

For a long time nobody took advantage of that fact; Not only different programming techniques were used to implement the interaction between the agent and the environment but also different programming languages and different platforms were used. This created a natural level of complexity which prevented the exchange of learning algorithms and test problems.

Eventually Tanner and White (2009)[17] created an open source software called RL-Glue. This software lays a foundation for empirical studies in reinforcement learning because it provides a necessary set of standards which allow to connect the test problems with the learning algorithms from different researchers without creation of this mundane connecting code. Essentially, RL-Glue, by introducing a set of standard functions, allows to plug-in agents, environments and experiments from different authors even created in different languages.

RL-Glue consists of the three main modules: an agent, an environment and an experiment. The agent is the module which encapsulates the learning algorithm. The environment is the module which encapsulates a particular test problem. The experiment is the program which serves as traditional main function and is responsible for running

the agent through the environment and collecting all the data obtained by the agent. The experiment module is place where researcher controls how many episodes or how many trials should be performed by the agent. All the interaction between the modules is handled by a central server.

The plug-in nature of the the RL-glue perfectly addresses three aspects of adequate

How RL-Glue Interacts with the Experiment Program, Agent and Environment

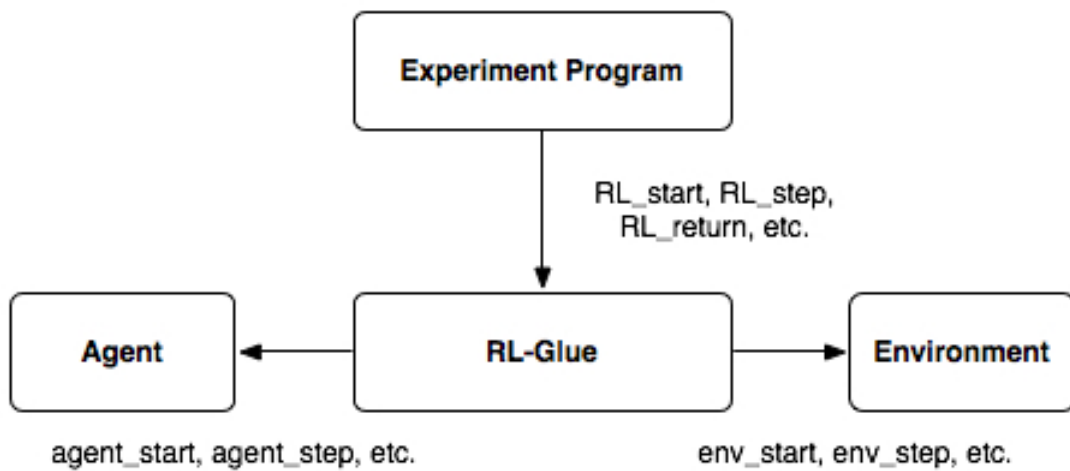


Figure 2.6: The RL-Glue standard.

methodology that were outlined before. Namely:

- The agent module encapsulates the learning algorithm to be evaluated. Each learning algorithm is a separate plug-in. Multiple number of the learning algorithms could be attached to the RL-glue simultaneously. This allows to evaluate multiple number of the algorithms simultaneously. A researcher just chooses which of the connected agents will be evaluated.
- The environment module tackles aspect number two of the adequate methodology. Because the environment module is well-documented and well-standardized in terms of standard programming interfaces it allows to create standard test problems. This modules eliminates the need of connecting codes needed at the present moment.

With plug-in architecture the researcher only needs to download a test problem and plug it into running RL-glue server without any connecting code.

- The experiment module tackles aspect number three of the adequate methodology. The experiment module is the plug-in where the researcher chooses the particular test problem and the particular learning algorithms that will be evaluated against that test problem. The experiment module is also a place where all the statistics is gathered and analyzed. This is achieved by introducing a standard message protocol between the agent module and the experiment module. Using this message protocol the agents can send information about learned Q-functions and policies to the experiment module. This is also a place where different performance metrics are applied. Each performance metric is created as a library with external functions which encapsulates the inner routines. Such performance metrics are stored separately in the library files and also could be accessible to the public. For example Hypervolume performance metric will be introduced to RL-glue as a part of this thesis and will be available publicly. So if any other researcher wants to use it, the only job left to do is to download it and add to the existing libraries. When the performance metric library is downloaded and introduced to RL-glue it offers an interface of public functions which accept information obtained from the agent.

Originally, RL-Glue is a software for single objective reinforcement learning but it can be adapted to multi-objective(which is also a part of this thesis) by replacing scalar reward signal with vector-valued.

2.3 Benchmarks.

Each benchmark or test problem is created to address a particular aspects of real-life problems. Thus a rich library of benchmarks should be available to adequately assess a learning algorithm. A suite of these benchmarks will allow to look at the learning from

different point of views. Combined together, the benchmarks will allow to evaluate how different algorithms cope with following aspects of real-life problems:

- Multiple number of the objectives. To test which algorithms cope better when the number of objectives is increasing.
- A range of benchmarks should be devoted to assess algorithms under stochasticity of reward and state transition functions.
- One of the basic requirements is to create benchmarks with continuous state and action spaces.
- Another important class of benchmarks is ones describing real-life problems with large dimensionality of state or action spaces that require the use of function approximation;
- Problems with partially-observable states.
- A mixture of episodic and continuing tasks;
- Each real-life problem assumes existence of Pareto front of optimal solutions. This front exhibits a number of different characteristics. A number of benchmarks should be created to expose the learning algorithms to all possible characteristics of the Pareto front.

(Vamplew et al, [?]) already introduced a number of benchmarks with known Pareto fronts. This greatly facilitates in the evaluation of the multi-objectives algorithms.

2.3.1 Deep sea treasure.

This test problem is a grid world consisting 10 rows and 11 columns as shown on figure 3.5. The grid represents undersea surface with multiple number of treasure spots available. One small remark - is that each treasure spot may have different treasure value. The

agent is represented as a submarine. Four actions available to the agent - left, right, up, down. Each of the actions move submarine by one square in appropriate direction. First objective is pretty obvious - to find as much as possible of treasure. Second objective is to minimize time spent.




Figure 2.7: Grid world for Deep Sea Treasure problem.

This problem comes with a known Pareto front which is shown in Figure 3.6. The form of the front is globally concave with number of local concavities.




Figure 2.8: Pareto front for Deep Sea Treasure problem.

2.3.2 MO-Puddleworld.

This test problem represents 2 dimensional world with puddles located across the world as shown in Figure 3.7. Each episode agents starts at a random location and must reach the goal position(top-right corner) while avoiding the puddles. Four actions available to the agent - left, right, up, down. In this problem reward vector is represented as two penalties: first applied each time step until goal state is reached, second applied when the agent enters a puddle.




Figure 2.9: Grid world for the Puddleworld problem.

The Puddleworld problem also comes with a known Pareto front which is shown in Figure 3.8.

From first inspection it is seen that the shape of the front is convex but at a closer inspection reveals a number of concavities and linearities.

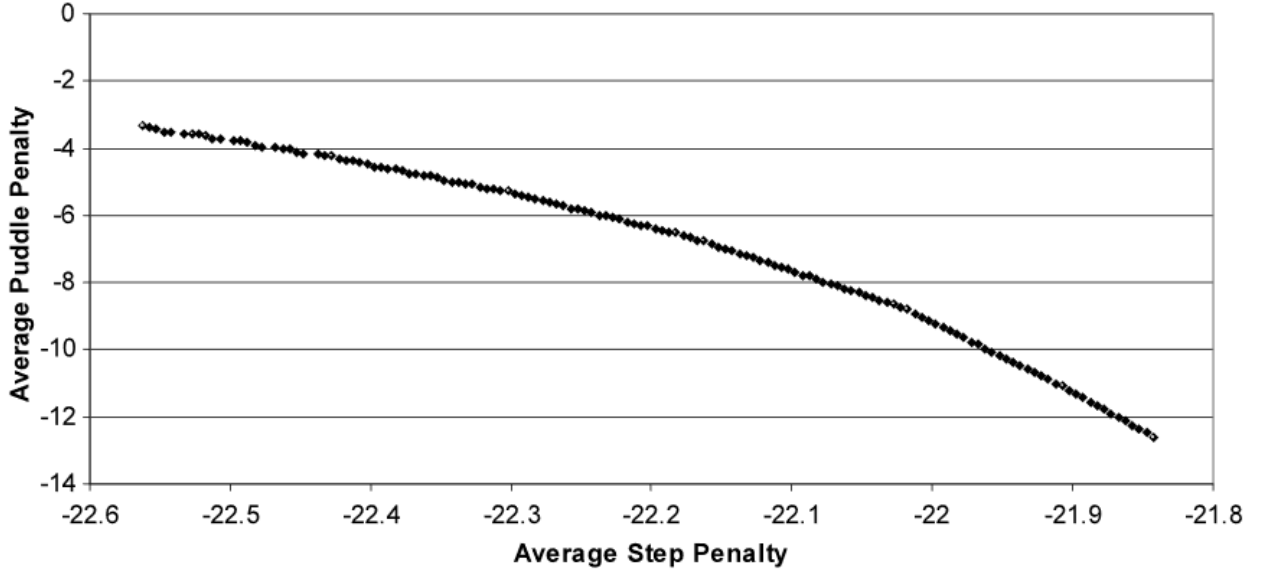


Figure 2.10: Pareto front for the Puddleworld problem.

2.3.3 MO-Mountain Car.

This test problem represents a valley with a car that must escape from the valley. But the car's engine is too weak to overcome the gravity so the car should climb left side of the valley with reverse acceleration to build enough potential energy to escape from the right side. Three actions available to the agent - full throttle forward, full throttle backward, and zero throttle. A penalty of -1 is received on all states except the goal state. There are two additional penalties: -1 for each reverse acceleration and -1 for each forward acceleration. So totally reward vector consists of three elements.

2.3.4 Resource gathering.

This test problem represents 2 dimensional grid world. Two types of resources are available for gathering - gold and gems. The goal of the agent is to gather either of both of the resources. Also there are locations where the hero can be attacked, the chance of being attacked is 10 percent. If the attack happens then the agent loses all gathered resources and returns to base. Four actions available to the agent - left, right, up, down. Each of

the actions move the agent by one square in appropriate direction. The reward vector consists of three elements and is structured in the following way - [enemy, gold, gems].




Figure 2.11: Grid world for the Resource Gathering problem.

(Barrett, Narayanan, [?]) used Convex Hull Value Iteration algorithm and found six optimal policies as shown on Figure 3.10.




Figure 2.12: Optimal Policies for the Resource Gathering problem.

References

- [1] Aissani, N., Beldjilali, B. and Trentesaux, D., *Use of machine learning for continuous improvement of the real time heterarchical manufacturing control system performances*, International Journal of Industrial and Systems Engineering, 2008, vol. 3(4):pp. 474–497
- [2] Barrett, L. and Narayanan, S., *Learning all optimal policies with multiple criteria*, in *Proceedings of the 25th international conference on Machine learning*, ACM, 2008 pp. 41–47
- [3] Bellman, R., *Dynamic programming and lagrange multipliers*, Proceedings of the National Academy of Sciences of the United States of America, 1956, vol. 42(10):p. 767
- [4] Berry, A. M., *Escaping the bounds of generality - unbounded bi-objective optimisation*, Ph.D. thesis, University of Tasmania, 2008
- [5] Bertsekas, D., *Dynamic programming and optimal control*, vol. 1, Athena Scientific Belmont, 1995
- [6] Castelletti, A., Corani, G., Rizzolli, A., Soncinie-Sessa, R. and Weber, E., *Reinforcement learning in the operational management of a water system*, in *IFAC Workshop on Modeling and Control in Environmental Issues, Keio University, Yokohama, Japan*, 2002 pp. 325–330

- [7] Feinberg, E. A., Feinberg, E. A., Shwartz, A. and Shwartz, A., *Constrained markov decision models with weighted discounted rewards*, Math. of Operations Research, 1993, vol. 20:pp. 302–320
- [8] Ferreira, L., Bianchi, R. and Ribeiro, C. H. C., *Multi-agent multi-objective learning using heuristically accelerated reinforcement learning*, in *Robotics Symposium and Latin American Robotics Symposium (SBR-LARS), 2012 Brazilian*, 2012 pp. 14–20, doi:10.1109/SBR-LARS.2012.10
- [9] Gábor, Z., Kalmár, Z. and Szepesvári, C., *Multi-criteria reinforcement learning.*, in *ICML*, vol. 98, 1998 pp. 197–205
- [10] Karlsson, J. and Ballard, D. H., *Learning to solve multiple goals*, 1997
- [11] Lizotte, D. J., Bowling, M. H. and Murphy, S. A., *Efficient reinforcement learning with multiple reward functions for randomized controlled trial analysis*, in *Proceedings of the 27th International Conference on Machine Learning (ICML-10)*, 2010 pp. 695–702
- [12] Natarajan, S. and Tadepalli, P., *Dynamic preferences in multi-criteria reinforcement learning*, in *Proceedings of the 22nd international conference on Machine learning*, ACM, 2005 pp. 601–608
- [13] Pareto, V., *Cours d\’economie politique*, 1896
- [14] Roijers, D. M., Vamplew, P., Whiteson, S. and Dazeley, R., *A survey of multi-objective sequential decision-making*, Journal of Artificial Intelligence Research, 2013, vol. 48:pp. 67–113
- [15] Shabani, N., *Incorporating flood control rule curves of the columbia river hydroelectric system in a multireservoir reinforcement learning optimization model*, Ph.D. thesis, University of British Columbia, 2009

- [16] Sutton, R. S. and Barto, A. G., *Time-derivative models of pavlovian reinforcement*, in *Learning and Computational Neuroscience: Foundations of Adaptive Networks*, MIT Press, 1990 pp. 497–537
- [17] Tanner, B. and White, A., *Rl-glue: Language-independent software for reinforcement-learning experiments*, The Journal of Machine Learning Research, 2009, vol. 10:pp. 2133–2136
- [18] Vamplew, P., Dazeley, R., Berry, A., Issabekov, R. and Dekker, E., *Empirical evaluation methods for multiobjective reinforcement learning algorithms*, Machine learning, 2011, pp. 1–30
- [19] Vamplew, P., Yearwood, J., Dazeley, R. and Berry, A., *On the limitations of scalarisation for multi-objective reinforcement learning of pareto fronts*, AI 2008: Advances in Artificial Intelligence, 2008, pp. 372–378
- [20] Van Moffaert, K. and Drugan, A., Madalina M, *Scalarized multi-objective reinforcement learning: Novel design techniques*, IEEE SSCI, 2012, vol. 13:pp. 94–103
- [21] Van Moffaert, K., Drugan, M. M. and Nowé, A., *Hypervolume-based multi-objective reinforcement learning*, in *Evolutionary Multi-Criterion Optimization*, pp. 352–366, Springer Berlin Heidelberg, 2013
- [22] Watkins, C. J. and Dayan, P., *Q-learning*, Machine learning, 1992, vol. 8(3-4):pp. 279–292
- [23] Yaman, F., Walsh, T. J., Littman, M. L. and desJardins, M., *Democratic approximation of lexicographic preference models*, Artificial Intelligence, 2011, vol. 175(78):pp. 1290 – 1307, ISSN 0004-3702, doi:<http://dx.doi.org/10.1016/j.artint.2010.11.012>, URL <http://www.sciencedirect.com/science/article/pii/S0004370210002018>, representing, Processing, and Learning Preferences: Theoretical and Practical Challenges

- [24] Zitzler, E. and Thiele, L., *Multiobjective evolutionary algorithms: a comparative case study and the strength pareto approach*, Evolutionary Computation, IEEE Transactions on, 1999, vol. 3(4):pp. 257–271, ISSN 1089-778X, doi:10.1109/4235.797969